



Workshop AI su Databricks Global Markets Trading & Investments

INTESA  SANPAOLO

Milano, 19 Maggio 2026



Agenda

15:30 – 15:40 | Apertura e obiettivi del workshop

15:40 – 15:50 | AI in Intesa Sanpaolo

15:50 – 16:00 | Databricks come piattaforma GenAI per Capital Markets

16:00 – 16:30 | DEMO 1 – Analytics in Capital Markets con Databricks AI/BI

16:30 – 16:55 | DEMO 2 – Databricks Coding

16:55 – 17:20 | DEMO 3 – GenAI per Market Intelligence e Research Augmentation

17:20 – 17:40 | GenAI, rischio e compliance nel trading

17:40 – 17:50 | Q&A, sintesi e prossimi passi



L'AI per il desk di Trading



Alcuni pain ricorrenti per i trading desk 1/2

EXECUTION Information overload & Decision Fatigue	MANIFESTAZIONE OPERATIVA I trader si devono tenere informati operando in un contesto di notizie e notifiche numerose ed eterogenee.	IMPATTO SISTEMICO Tempo e attenzione consumati dall'aggregazione delle notizie e dalla valutazione delle notifiche. [1]	LIMITE DEI PROCESSI News provider volti al coverage informativo, non alla velocità decisionale. Alert si basano su soglie, non sono contestuali al book.	STATO TARGET Flusso di notifiche classificato per natura, book, urgenza e impatto. Alert spiegabili, pochi, con priorità esplicita.
EXECUTION Dark data e segnale non strutturato	MANIFESTAZIONE OPERATIVA Il trader opera in contesti di dati non strutturati che influenzano le decisioni in modo implicito.	IMPATTO SISTEMICO Segnali qualitativi (market colour o client intent) non diventano sistematici. Apprendimento organizzativo debole.[2]	LIMITE DEI PROCESSI Formati eterogenei senza layer di strutturazione unificato. Mancanza di strutture per leggere dati in formati non tradizionali.	STATO TARGET Sistema di interrogazione in linguaggio naturale di report e comunicati, research, chat, voice, report e comunicazioni. Segnali qualitativi in output interrogabili e archiviabili.
WORKFLOW Data silos & Frammentazione	MANIFESTAZIONE OPERATIVA La banca ha la posizione nel front office system, nel risk system un'esposizione aggregata su feed ritardato, nel back office la posizione contabile post-settlement.	IMPATTO SISTEMICO Tempo speso e carico cognitivo per riconciliazioni tra sistemi. Carico organizzativo invisibile e distribuito che sottrae capacità decisionale al desk.	LIMITE DEI PROCESSI Copie della stessa informazione in sistemi strutturalmente differenti per timing e viste (EOD, EOM, intraday).	STATO TARGET Layer di orchestrazione su tutti i sistemi con rilevazione automatica di discrepanze e attribuzione. Investigazione manuale solo per eccezioni reali.

Alcuni pain ricorrenti per i trading desk 2/2

Risk Monitoring EXECUTION	MANIFESTAZIONE OPERATIVA I breach di risk sono tardivi poiché basati su dati EOD. Il desk viene notificato a breach avvenuto, specialmente i limiti legati a concentrazione o stress scenario.	IMPATTO SISTEMICO Il trader rischia di eseguire senza piena visibilità del contributo marginale al rischio del book, basandosi su pre-trade check che puntano a parametri non aggiornati.	LIMITE DEI PROCESSI Calcoli batch e stress test overnight: non riflettono l'andamento dell'intraday. Monitoraggio progettato per compliance normativa, non per la gestione operativa.[20]	STATO TARGET Potenza computazionale adeguata per exposure, sensitivity e P&L risk-adjusted quasi in real-time. Alert su correlazioni emergenti, non solo su limiti statici. Contributo marginale visibile prima dell'execution.[1]
Ricerca Quant EXECUTION	MANIFESTAZIONE OPERATIVA Una nuova ipotesi richiede giorni: richiesta dati, pulizia, feature engineering, backtest, review IT. Il collo di bottiglia è l'infrastruttura. Ogni iterazione lenta riduce la possibilità di esplorare alternative.	IMPATTO SISTEMICO Si rischia di inficiare l'Alpha discovery: meno ipotesi testate, minore qualità del modello. Solo il 19% delle organizzazioni AI raggiunge il deploy on-scale. [4]	LIMITE DEI PROCESSI Catalogo dati assente o limitato. Storage storico separato dal flusso live. Backtesting seriale dove serve parallelismo. Ricerca non riproducibile per mancanza di versioning. Pipeline manuali su ogni deploy.	STATO TARGET Piattaforma unica per accesso self-service ai dati storici e live. Compute distribuito per backtest e parameter sweep. Feature store condiviso. Pipeline continua da ricerca a paper trading. Governance integrata: 12× più progetti in produzione.[4][6]
Reporting verso l'esterno WORKFLOW	MANIFESTAZIONE OPERATIVA Report verso responsabili, funzioni di controllo o clienti: processo ripetitivo di gestione manuale, verifica e formato. Dipendenza da IT.	IMPATTO SISTEMICO Tempo speso per lavoro a basso valore aggiunto. Stakeholder ricevono dati non aggiornati. Eventuali errori nei report comportano ricicli.	LIMITE DEI PROCESSI Nessuna integrazione tra sistemi operativi e strumenti di reporting. La narrativa di commento è sempre manuale.	STATO TARGET Generazione automatica di report certificati da dati governati. Draft narrativo AI abbinato a revisione umana. Report regolamentari da dati Unity Catalog.[3][5]

Cosa fa oggi l'AI per chi lavora sul desk?

WORKFLOW & PRODUCTIVITY

70%

dei capital markets firms intervistate da Databricks ha investment analytics come modalità d'uso AI principale [7][8]

Linguaggio naturale

Il trader interroga i dati in linguaggio naturale, senza dipendere dall'IT. Goldman Sachs ha rilasciato il GS AI Assistant come tool di accesso a dati e knowledge interna per tutti i dipendenti.[11]

MARKET INTELLIGENCE

Market intelligence

AI estrae intelligence da fonti pubbliche ed interne sostituendo lo screening manuale con segnali real-time e brief sintetici tracciabili. Il fondo Two Sigma gestisce \$60B+ AuM con strategie AI-driven.[7]

Predictive analytics

Modelli ML su serie storiche di volatilità, liquidità e prezzi: pipeline quant scalabili documentate da Databricks per backtesting su tick data di FX, equity e fixed income.[10]

EXECUTION & RISK

Ottimizzazione book

Scenario analysis e stress test del book su migliaia di simulazioni in secondi. LSEG e Databricks hanno integrato tick history e Agent Bricks per ottimizzazione di portafoglio in capital markets. [9]

Risk Monitoring

Esposizioni e correlazioni in real-time vs snapshot EOD. JPMorgan rivende "Risk as a Service" (la tecnologia dei propri desk).[22] Fed TCAM e IMF GFSR suggeriscono l'intraday monitoring.[20][21]

What-if pre-trade

L'impatto di un'operazione su P&L e rischio del book calcolato in secondi. Databricks MCP abilita intraday repricing, stress scenarios e hedging analysis in real-time.[6]



Piattaforma MLOps / AI di Gruppo e Servizi trasversali






Databricks come piattaforma GenAI per Capital Markets



Our learnings from 600+ Capital Markets players




Industrialized AI is driving measurable margin expansion

Leading firms have shifted **from isolated pilots to enterprise-scale AI strategies** that automate high-value workflows to prioritize ROI over experimentation



Modernized data estates are the prerequisite for Agentic AI

Success hinges on a **unified, governed** data architecture that eliminates legacy **silos** to fuel "agentic" systems capable of automating complex, multi-step market processes



Proactive oversight is moving firm-wide control to machine-speed

The industry has moved **from static**, report-based monitoring **to real-time "intelligence loops"** that integrate AI into core data flows to detect and triage incidents in minutes




Information Edge depends on the extended enterprise ecosystem

Differentiation no longer comes from internal data alone, but from using **open sharing protocols** to seamlessly ingest and monetize insights across a global partner network



Our mission is to democratize Data & AI



**Ingest & aggregate
data across all
sources**

**One copy of data
at the lowest
TCO—foundational
for AI-readiness.**




**Secure & govern
the Data & AI
"brain"**

**Secure, monitor,
audit, discover and
share data and AI
assets across the
enterprise.**



**Democratize data
across enterprise**

**Access to all data
across all users, via
conversational
interfaces.**



**Build domain-
specific AI**

**Develop agentic &
machine learning
systems,
leveraging open
ecosystem
innovation.**

Data Intelligence Platform

The Data and AI Intelligence Platform



Agents

Production AI agents



AI/BI

Agentic business intelligence



Custom Apps

Secure data and AI apps

ML and more....



AI with Enterprise Context



Genie



Lakehouse

Data warehousing



Lakebase

Serverless Postgres



Lakeflow

Ingest, ETL, streaming



Unified Governance



Unity Catalog



Open Formats



Postgres



DELTA LAKE

ICEBERG



Databricks is entrenched in the capital markets ecosystem

MARKET DATA & REFERENCE

Market Data

FACTSET

S&P Global

Bloomberg



LSEG

Compliance

MOODY'S

Exchanges



Nasdaq

Alternative & Specialty Data



Stocktwits

yipitDATA

Nasdaq

M Science

MIDDLEWARE & INTEGRATION

Applications & Orchestrations

OMEGA POINT



Outcome Catalyst

KX

Custom Apps
Secure data and AI apps

AI/BI
Agentic business intelligence

And more...



Lakeflow
Ingest, ETL, streaming



DBSQL
Data warehousing



Lakebase
Serverless Postgres



Agent Bricks
Agentic AI & Machine Learning

Unity Catalog



DELTA LAKE



ICEBERG



Postgres

System Integrators

accenture

Deloitte.

EY

Capgemini



IMPETUS

slalom

rearc

TREDENCE

TIGER
ANALYTICS

Datagum
Data solutions that stick

[DATA
INTELLECT]

FRONT OFFICE

Master Data Management

FACTSET



LakeFusion

AI Agents & Automation

FACTSET

S&P Global



LSEG

Nasdaq

ANTHROPIC

Client Solutions & Servicing

ADDEPAR

rimes



The capital markets industry runs on Databricks



600+

capital markets
players globally

2,000+

financial
services clients

500+

data listings on
marketplace

250

ISV and
SI partners



Why **capital** **markets** choose Databricks



Open by design

Zero lock-in with open formats, multi-cloud flexibility and full interoperability across systems and teams



Built for the markets

A secure foundation with unified governance and end-to-end lineage, ensuring data is instantly auditable and compliant with complex regulations



Single platform

Connects your data into AI-powered value from “talk to your data” with Genie to real workflow apps end to end



Lower TCO

AI-optimized throughout and gets more efficient as it learns your data, saving you money





RBC Capital Markets: Transforming research with AI to accelerate turnaround times by 60%

Challenge

RBC Capital Markets faced mounting pressure to deliver **expert research and insights** across a broader coverage universe, faster than ever before, as clients demanded **near real-time**, differentiated commentary.

Traditional analyst workflows were **highly manual** and **slow**, with post-earnings research notes taking up to 45 minutes per company—creating bottlenecks and risking missed market opportunities.

Solution

RBC partnered with Databricks to build Aiden QuickTakes, a proprietary AI application embedded within a secure, cloud-native **Data Intelligence Platform**.

This unified environment—leveraging Databricks' **Unity Catalog** for governance and **Mosaic AI** for end-to-end AI capabilities—automates data ingestion, analysis, and content generation, enabling analysts to produce high-quality research drafts in minutes while maintaining strict security and compliance standards.

Impact

- ▲ **60% Faster Turnaround:** Reduced the end-to-end research process from 45 minutes down to just 15 minutes with "human-in-the-loop" validation.
- ▲ **Scaled Global Coverage:** Reclaimed thousands of hours, allowing analysts to expand company coverage and focus on high-value client dialogue rather than manual data entry.





State Street: Slashes response times by 99% and prevents financial loss with real-time RiskOps

Challenge

State Street managed \$47T in assets but struggled with a reactive risk posture. Legacy, **report-based systems** could not analyze data embedded in financial flows in real-time.

This prevented the firm from acting on market volatility and sophisticated cyber threats at machine speed, creating significant financial and regulatory risk.

Solution

Built an AI-powered platform on Databricks using a "**Risk Intelligence Loop**" (Detect, Triage, Act, Validate) to move from static reporting to continuous intelligence.

Leveraged the **Lakehouse** to treat data as a product and utilized **MLFlow** to retrain models and rewire workflows in a continuous, virtuous cycle.

Impact

- ▲ **Slashed Cycle Times:** Complex reconciliation dropped from 2-3 days to < 4 hours, while human decision cycles fell to < 30 minutes.
- ▲ **Machine-Speed Precision:** Predictive break detection accuracy jumped to > 80%, with 97%+ of resolutions now machine-driven, preventing incidents from becoming regulatory headlines.



NYSE/ICE: Transforming financial discovery with an AI-powered text-to-SQL application

Challenge

The Intercontinental Exchange (ICE) faced the challenge of **enabling non-tech users** to access structured data without needing to understand complex data models, schemas, or **SQL queries**.

This hindered their ability to efficiently extract insights from large volumes of financial data.

Solution

Pioneered a structured RAG application on the Databricks **Data Intelligence Platform** that translated natural language into SQL.

Leveraged **Mosaic AI** (Vector Search, Model Serving) and Unity Catalog to provide a text-to-SQL app that enabled users to generate accurate queries, enhancing data accessibility and operational efficiency.

Impact

- ▲ **Remarkable Accuracy:** Achieved a 96% execution match and 77% syntactic accuracy across test queries, ensuring reliable financial insights for non-technical users.
- ▲ **Rapid Time-to-Value:** Developed and deployed the robust text-to-SQL system from concept to production in just five weeks.





DEMO 1 – Analytics in Capital Markets con Databricks AI/BI




BI – AI/BI Dashboards

A modern, fast, and scalable BI solution for the entire organization


AI/BI offers rich visuals, cross-filtering, parameters, and high-performance client-side filters:

- The 24-hour shared cache enables near-instant load times for viewers.
- It does not require end-user licenses, making it ideal for scaling BI across the entire company.
- It is particularly effective for ad hoc analysis and large datasets where latency and cost matter.



Data Warehouse – front-end with AI/BI e Genie

From static dashboards to conversational self-service




AI/BI delivers very fast dashboards, and you only pay for query compute:

- Adoption grows quickly thanks to a dedicated UI for business users.
- Each dashboard can have a dedicated Genie for natural language questions.
- The cycle from new question to new BI view is drastically reduced.

BI – Genie for Business Users

Natural language analytics with enterprise context



Genie enables non-technical users to query data in natural language:

- It understands business semantics, KPIs, and joins defined in UC and in the space setup.
- It maintains conversational context, making follow-up questions more natural.
- Golden SQL, benchmarking, and monitoring help data teams improve its accuracy over time.





DEMO 2 – Databricks Coding







DEMO 3 – GenAI per Market Intelligence e Research Augmentation



The point:

Latency

In fast-moving markets, spreads can widen 10–15 basis points in hours. The opportunity to reduce exposure at favorable prices may have already passed.

Limited Coverage

Sarah must prioritize and potentially miss important signals. Subtle signals may go unnoticed until it's too late.

Inconsistent Methodology

Different historical comparisons across PMs, unable to maintain consistent risk standards.

Data Fragmentation

OMS, excel for calculation, LSEG application for pricing, all that context switch adds time and creates opportunities for errors.

No Audit Trail

There's no systematic record of the decision-making process. "Why did you assume 15 bps of widening?"

Agentic Bond Trading Monitoring

Combines AI reasoning, real-time data integration, autonomous execution and human in the loop with **Agent Bricks**

News breaks at 7:30 AM

System scrapes article automatically from Delta Share

7:31 AM

LLM extracts affected sectors and key facts (Agent Bricks)

7:32 AM

Database queries portfolio exposure across sectors (Genie)

7:34 AM

LLM search for historical shocks (Vector Store)

7:37 AM

Model predict 30 days spread movements (Classical ML)

7:38 AM

LSEG repricing calculates exacts P&L impact (External MCP)

7:40 AM

Complete report with recommendations (Foundation Model)

7:47 AM

Sarah arrives at her desk with her usual large coffee and finds a complete analysis waiting in her inbox. The decision has already been made, documented, and quantified. She reviews the recommendations, validates the logic, and approves the trades.

8:00 AM

By market opening, the team has reduced exposure at favorable prices, mitigated \$85k in estimated losses, and moved on to the next opportunity.

This level of automation can only be achieved by a platform capable of...

Handling both structured and unstructured data (filings, news)

Supporting different foundation models and agents through MCP server integration

Governing all data assets to keep track of context feeding models for explainability and transparency



AI/BI Genie

Turn your tables into an expert AI chatbot.

Text-to-SQL
Table Q&A
Structured data Q&A

Open Genie

Information Extraction

Extract key information & insights into structured JSON.

Classification
Sentiment analysis
Entity resolution
Extraction
Summarization

Build

Knowledge Assistant

Turn your docs into an expert AI chatbot.

RAG
Document Q&A

Build

External MCP Server

Connect functions or other agents


API
External MCP

Multi-Agent Supervisor

Design an AI system with Genie, agents, tools.

Tool-calling
Agent orchestration

Build



Using Databricks One, Sarah combines 3rd party and 1st party data for portfolio management through Genie, running portfolio analysis in natural language.

 AI/BI Genie

Turn your tables into an expert AI chatbot.

Text-to-SQL Table Q&A

Structured data Q&A

Open Genie 

 Information Extraction

Extract key information & insights into structured JSON.

Classification Sentiment analysis

Entity resolution Extraction Summarization

Build

 Knowledge Assistant

Turn your docs into an expert AI chatbot.

RAG Document Q&A

Build

 External MCP Server

Connect functions or other agents


API External MCP

 Multi-Agent Supervisor

Design an AI system with Genie, agents, tools.

Tool-calling Agent orchestration

Build



With Reinforcement Learning through Human Feedback, Agent Bricks optimizes task of extracting company information from corporate news events. Symbols, company names, industry sectors and key topics.

AI/BI Genie

Turn your tables into an expert AI chatbot.

Text-to-SQL Table Q&A

Structured data Q&A

Open Genie

Information Extraction

Extract key information & insights into structured JSON.

Classification Sentiment analysis

Entity resolution Extraction Summarization

Build

Knowledge Assistant

Turn your docs into an expert AI chatbot.

RAG Document Q&A

Build

External MCP Server

Connect functions or other agents


API External MCP

Multi-Agent Supervisor

Design an AI system with Genie, agents, tools.

Tool-calling Agent orchestration

Build



Indexing news events for semantic search and RAG application, searching for similar events / trends in the past as input shocks to credit spread predictive models

 **AI/BI Genie**

Turn your tables into an expert AI chatbot.

Text-to-SQL Table Q&A
Structured data Q&A

[Open Genie](#)

 **Information Extraction**

Extract key information & insights into structured JSON.

Classification Sentiment analysis
Entity resolution Extraction Summarization

[Build](#)

 **Knowledge Assistant**

Turn your docs into an expert AI chatbot.

RAG Document Q&A

[Build](#)

 **External MCP Server**

Connect functions or other agents


API External MCP

 **Multi-Agent Supervisor**

Design an AI system with Genie, agents, tools.

Tool-calling Agent orchestration

[Build](#)



Interoperability of platforms & systems through MCP server, connecting both internal and external applications for risk calculations / pricing engines


 **AI/BI Genie**
Turn your tables into an expert AI chatbot.
Text-to-SQL Table Q&A
Structured data Q&A
[Open Genie](#)

 **Information Extraction**
Extract key information & insights into structured JSON.
Classification Sentiment analysis
Entity resolution Extraction Summarization
[Build](#)

 **Knowledge Assistant**
Turn your docs into an expert AI chatbot.
RAG Document Q&A
[Build](#)

 **External MCP Server**
Connect functions or other agents
API External MCP

 **Multi-Agent Supervisor**
Design an AI system with Genie, agents, tools.
Tool-calling Agent orchestration
[Build](#)



AI system runs 24x7, monitoring global news flows (CDC). The moment a material event breaks, it instantly identifies affected positions, calculates portfolio exposure, predicts price movements, reprices holdings and delivers specific trade recommendations.



GenAI, rischio e compliance nel trading



Governance AI: quattro aree non negoziabili

01 Human-in-the-loop

Ogni output AI rilevante passa sotto revisione umana prima di essere azionato. Il sistema AI suggerisce; il trader decide. Nessun ordine AI-generated senza approvazione esplicita.[12][13]

02 Tracciabilità di dati, prompt e output

Ogni interazione AI registrata: input, modello (versione e parametri), output, utente, timestamp. Tracciabilità completa per audit regolamentare e investigazione interna.[5][12]

03 Mitigazione rischi AI

Hallucination, bias e prompt injection sono rischi operativi reali nel trading. RAG, guardrail input/output, red teaming e drift monitoring mitigano sistematicamente i rischi.[14][16]

04 Allineamento normativo

Ogni sistema AI in produzione richiede conformity assessment, model card e integrazione nel framework di model risk.

- EU AI Act (high-risk classification),
- MiFID II Art.17/48 (algorithmic trading controls),
- DORA

Controllo umano e tracciabilità delle decisioni AI

Human-in-the-loop e responsabilità decisionale	PRINCIPIO OPERATIVO Il trader rimane sempre responsabile della decisione finale. L'AI suggerisce, priorizza, allerta — ma non agisce autonomamente su decisioni material. Approval workflow obbligatorio prima dell'execution. Override sempre disponibile.	RISCHIO SE IGNORATO Responsabilità legale ambigua su AI-generated trades. Impossibile dimostrare oversight a supervisori e audit. Violazione EU AI Act Art.14 (high-risk AI systems). Risk di sanzioni regolamentari e perdita di fiducia degli investitori.	COME SI IMPLEMENTA Step di approvazione umana pre-execution nel workflow. Log immutabile di chi ha approvato, quando e perché. Escalation automatica su soglie di rischio. Dashboard di monitoring delle decisioni AI-assisted in tempo reale.	NORMA DI RIFERIMENTO • EU AI Act Art.14 — Human oversight[12] • MiFID II Art.17, Art.48 — Algo trading controls[17] • ESMA AI in securities markets[13] • Basel BCBS d530 — Risk governance[18]
Auditabilità di dati, prompt e output	PRINCIPIO OPERATIVO Ogni interazione AI registrata in modo strutturato: dati di input, prompt esatto, versione del modello, parametri, output generato, utente, timestamp. Tracciabilità completa e immutabile su Unity Catalog.	RISCHIO SE IGNORATO Impossibile rispondere a richieste di audit MiFIR/DORA. Zero evidence per dimostrare compliance EU AI Act. Nessuna visibilità su come le decisioni AI-assisted sono state prese. Rischio di sanzioni fino a €30M.	COME SI IMPLEMENTA Lineage automatico su ogni dato, prompt, modello e output. MLflow: model versioning e audit trail completo. Logging strutturato nativo. Accesso controllato per ruolo.	NORMA DI RIFERIMENTO • EU AI Act Art.17-20 — Record-keeping[12] • MiFIR Art.25 — Obligation to keep records[5] • DORA Art.10 — ICT incident log[15] • ESMA guidelines on AI in markets[13]

[5] MiFIR Reg.600/2014 · [12] EU AI Act Reg.2024/1689 · [13] esma.europa.eu · [15] DORA Reg.2022/2554 · [17] MiFID II Dir.2014/65/EU · [18] bis.org/bcbs/publ/d530.pdf

Guardrail tecnici e framework regolamentare

Mitigazione rischi (hallucination, bias, misuse)	PRINCIPIO OPERATIVO LLM può generare dati plausibili ma falsi, citazioni normative errate o prezzi inventati (<i>hallucination</i>). Modelli addestrati su dati storici possono amplificare trend passati e sottostimare tail risk (<i>bias</i>). Prompt injection da fonti esterne è un rischio operativo reale(misuse).	RISCHIO SE IGNORATO Trade basati su dati inventati dall'AI. Allocation distorta da bias demografici o di mercato. Manipolazione del modello via prompt injection da research esterno o email. Perdite dirette e rischio reputazionale.	COME SI IMPLEMENTA RAG (Retrieval-Augmented Generation) àncora le risposte a dati verificati. Guardrail input/output per bloccare contenuti non conformi. Red teaming periodico. Drift monitoring tra distribuzione di training e inference live.	NORMA DI RIFERIMENTO • NIST AI RMF 1.0 Govern/Measure/Manage[14] • FCA DP5/22 — AI/ML in financial services[16] • ESMA TechPaper AI in EU securities markets[13]
Allineamento con policy interne e regolamentazione	PRINCIPIO OPERATIVO Ogni sistema AI in produzione deve essere mappato rispetto alla classificazione EU AI Act (high-risk o GPAI), documentato per conformità, e integrato nel model risk management esistente. Il trading desk è l'utente finale ma non l'unico responsabile della governance.	RISCHIO SE IGNORATO EU AI Act: sanzioni fino a €30M o 6% del fatturato globale per sistemi high-risk non conformi.[12] DORA: incident reporting e resilience testing obbligatori. Gap competitivo vs concorrenti già in compliance strutturata.	COME SI IMPLEMENTA Model card per ogni modello AI in produzione. Conformity assessment per sistemi high-risk. Allineamento policy interne model risk.	NORMA DI RIFERIMENTO • EU AI Act Reg.2024/1689 — High-risk AI[12] • DORA Reg.2022/2554 — Digital resilience[15] • BCBS d562 — Operational resilience[19] • MiFID II Art.17/48 — Algo controls[17]

[12] EU AI Act · [13] esma.europa.eu · [14] NIST · [15] DORA · [16] FCA DP5/22 · [17] MiFID II · [19] BCBS d562

Q&A, sintesi e prossimi passi




FSI Forum Milano

Il punto d'incontro per i professionisti dei dati e dell'IA nei settori Banking, Capital Markets, Audit e Insurance in Italia

Mercoledì , 27 maggio, 2026 | 9:00 – 14:00

UNA Hotel Milano Verticale



NEXT STEPS

- Discussione sui use case più rilevanti per la struttura
- Possibili next step:
 - proof of concept su un desk
 - estensione ad altre asset class
 - roadmap di adozione



GRAZIE!



databricks